

Exercise 2

*Configuring Auto Scaling and Load Balancing
Deploying a node.js server connected to an RDS database.*

Prior Knowledge

Unix Command Line Shell
EC2 starting servers

Learning Objectives

How servers interconnect in EC2
Passing configuration and automating setup of services in EC2
AutoScaling (without load balancing)

Software Requirements

- AWS CLI

Part A: Starting an instance with a userdata configuration

1. In our previous lab, we installed and started Apache by hand in the EC2 instance. Obviously that is not a tenable approach for a real production system. There are several options that could replace this:
 - a. We could set up a server by hand and then save the configuration to a new AMI image and use that in future.
 - b. We could utilize Docker and containers (we'll talk more about this later)
 - c. We could use configuration management tools like Puppet, Chef, Salt or Ansible. Or Amazon's own OpsWorks (which uses Chef)
 - d. But those go beyond the scope of this class, so we are going to use a simpler approach based on Amazon's "userdata" which allows us to pass a startup script to the newly launched instance.
2. EC2 allows us to pass a script that is run as root. This is passed in a format called userdata.
3. Go back to the console (and login if you need to again)
<https://ox-clo.signin.aws.amazon.com/console>
4. There is already an Amazon Aurora (MySQL compatible) database running in the cloud. It has a small amount of data in it that we will query from a node.js application. If you go to the RDS section of the AWS management console you can take a look at this instance. Please do not modify it!
5. Now let's try the instance manually before we create an auto-scaling version.

6. Go to the EC2 console, and Launch a new instance.
7. Choose the **Ubuntu Server 18.04 LTS (HVM)**
8. Once again choose a t2.micro instance and then **Next: Configure Instance Details**
9. At the bottom of the page you will find a section called **Advanced Details**. Expand this.
10. This is where our script will go. We are going to paste it into the **User Data** section. (You could also create a file and upload that if you prefer)
11. In your browser go to <http://freo.me/oxclo-userdata>
12. Now copy and paste the startup script into the user data section. It looks like this:

```
#!/bin/bash
# verbosity
set -e -x
# update the package list
apt-get update
# install node, node package manager and git.
apt-get -y install nodejs npm git
# some node packages including forever expect nodejs to be called node
ln -s /usr/bin/nodejs /usr/local/bin/node
# use the node package manager to install express.js and mysql support
npm install express mysql
# forever is a daemon for running node.js code
npm install forever -g
# change to the ubuntu home directory
cd /home/ubuntu
# use git to copy the node.js code into the system
git clone https://github.com/pzfreo/auto-deploy-node-js.git
cd auto-deploy-node-js
# pass the DB connection parameters into the code
export DBURL=oxclo-cluster.citfamclcdxs.eu-west-1.rds.amazonaws.com
export DBUSER=node
export DBPW=node
# start the server as a daemon
forever start --minUptime=1000 --spinSleepTime=1000 clustertest.js
#that's all
```

The script is doing the following:

- i. Installing node.js, the node package manager (npm) and git
- ii. Using npm to install some node packages (mysql, express.js and forever)
- iii. Using git to install our source code:
<https://github.com/pzfreo/auto-deploy-node-js/blob/master/clustertest.js>

- iv. Setting up the URL, userid and password for the database.

Hint: If you are worried about the security of this password (which you should be) then consider this. Firstly, you would certainly not normally put this into a public git repository! Secondly, only instances in Amazon can connect to the database (I'll explain shortly). Thirdly, this userid/password only has the access rights to read a single table.

- v. Using the **forever** toolkit to run our node.js code

13. Click **Next: Add Storage**, then **Next: Add Tags**

14. Add the **Name** tag as before. This time use `<your userid>-node`.
E.g. `oxclo02-node`

15. Click **Next: Configure Security Group**

16. Select an existing security group in the dropdown menu

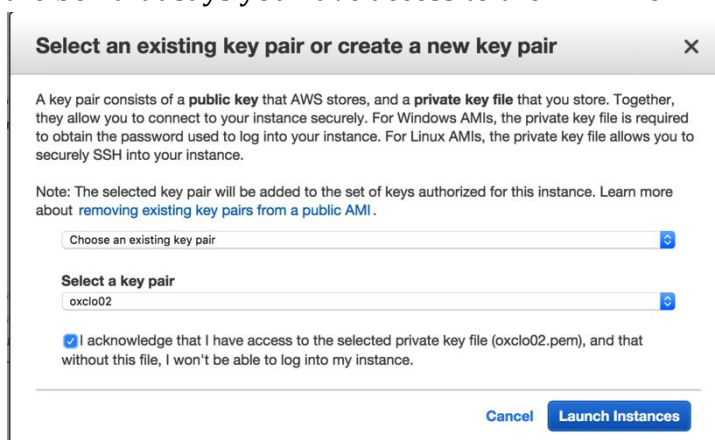
17. Choose the “node-security-group”

- a. *This is important because this group is allowed to access the database. We'll take a look shortly.*

18. Click **Review and Launch**

19. Click **Launch**

20. This time select to use an existing keypair, and find your own key pair. Check the box that says you have access to the PEM file:



Select an existing key pair or create a new key pair X

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Choose an existing key pair

Select a key pair

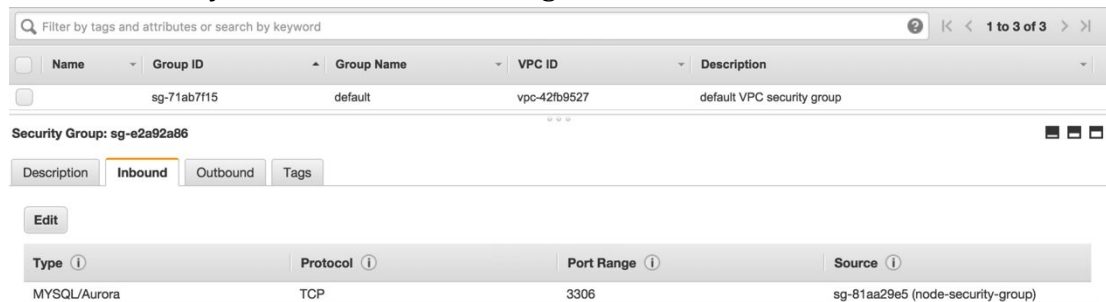
oxclo02

☒ I acknowledge that I have access to the selected private key file (oxclo02.pem), and that without this file, I won't be able to log into my instance.

Cancel Launch Instances

21. Now select **Launch Instances**

22. As before, go take a look at your instance status by clicking on the instance link.
23. While you wait for your instance to get going, you can take a look at the Security Groups. If you look at the **rds-security-group** and take a look at the inbound rules, you will see the following:

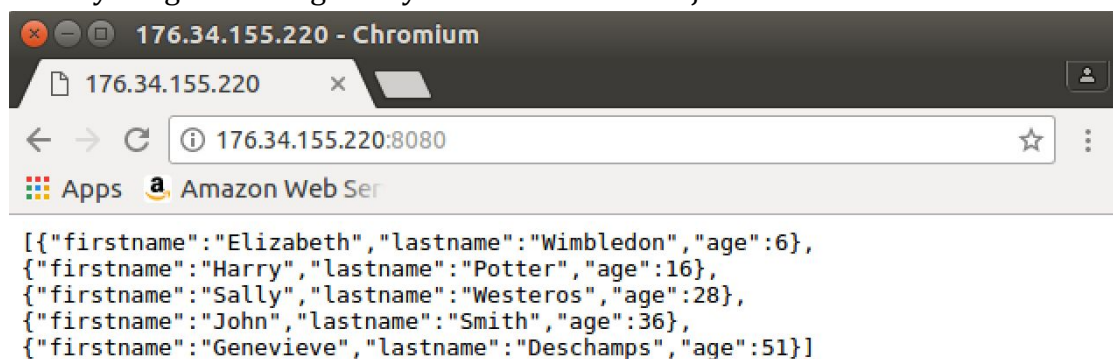


The screenshot shows the AWS Management Console interface for a Security Group. At the top, there's a search bar and a table with columns: Name, Group ID, Group Name, VPC ID, and Description. Below this, the 'Inbound' tab is selected for the security group 'sg-e2a92a86'. The 'Inbound' rules table shows a single rule with the following details:

Type	Protocol	Port Range	Source
MYSQL/Aurora	TCP	3306	sg-81aa29e5 (node-security-group)

What this shows is that only instances started in the node-security-group can access the RDS instances port 3306.

24. Go back and find your instance running (e.g. tagged `oxclo0n-node`).
25. If it has started and the status checks are finished, it may have completed its startup script. But this is a lot of work for a poor old micro instance to manage, so don't expect miracles.
26. Copy the public IP address of the instance and try browsing to <http://ww.xx.yy.zz:8080> (where the ww.xx.yy.zz are replaced with the public IP of your instance).
27. If everything is running then you should see some json returned.



28. If there is a problem, you can see the state of your startup by SSH-ing into your instance and doing:

```
tail -f /var/log/cloud-init-output.log
```

If this is still scrolling past then your server hasn't started up yet.

If this shows something like:

```
cloud-init v. 0.7.5 finished at Mon, 16 Jun 2015 22:30:26  
+0000. Datasource DataSourceEc2. Up 72.17 seconds
```

then the server init has started. Press Ctrl-C to exit tail.

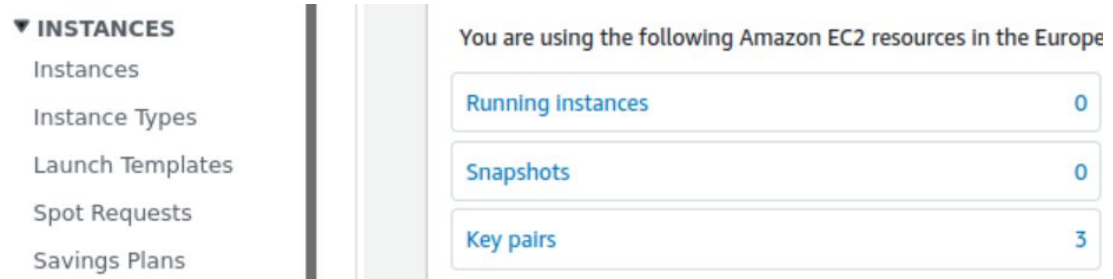
29. Once the server has started fully, try browsing again.

30. Once you have tested this, please **terminate** your instance through the AWS console.

PART B. Creating a Launch Template

31. In order to auto-scale this we are going to create a template for launching new servers. It is similar to creating a server but then we let Amazon decide when to start new servers.

32. In the **EC2 console**, you need to scroll the left hand menu to the bottom where you will find Launch Templates:



33. Click on **Launch Templates**

34. Now click on **Create Launch Template**

35. Use the name *userid-lt* (e.g. oxclo02-lt).

36. Give the template a description.

37. Choose AMI: Ubuntu Server 18.04 LTS (HVM), SSD

38. Choose Instance type t2.micro

39. So far your screen should look like:

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

Max 255 chars

Auto Scaling guidance [Info](#)
Select this if you intend to use this template with EC2 Auto Scaling

☐ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

► Template tags

► Source template

Launch template contents

Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.

Amazon machine image (AMI) [Info](#)

AMI

Ubuntu Server 18.04 LTS (HVM), SSD Volume Type
ami-089cc16f7f08c4457
Catalog: Quick Start architecture: 64-bit (x86) virtualization: hvm


▼

Instance type [Info](#)

Instance type

t2.micro
Family: General purpose 1 vCPU 1 GiB Memory
On-Demand Linux pricing: 0.0126 USD per Hour
On-Demand Windows pricing: 0.0172 USD per Hour

Free tier eligible ▼

[Instance types](#) 

40. Choose your key pair

41. Choose the **node-security-group**

42. Add an Instance Tag of:

Name: oxcloXX-scaled

43. Your next part of the screen should look like:

The screenshot displays the 'Create Launch Template' page in the AWS Management Console. It is divided into four main sections:

- Key pair (login)**: A dropdown menu for 'Key pair name' is set to 'oxclo01'. A 'Create new key pair' link with a download icon is to the right.
- Network settings**:
 - Networking platform**: Two options are shown. 'Virtual Private Cloud (VPC)' is selected, with a description: 'Launch into a virtual network in your own logically isolated area within the AWS cloud'. 'EC2-Classic' is unselected, with a description: 'Launch into a single flat network that you share with other customers'.
 - Security groups**: A dropdown menu shows 'Select security groups'. Below it, a selected group is displayed: 'node-security-group sg-20e61758' with a close icon, and 'VPC: vpc-42fb9527'.
- Storage (volumes)**: A section for 'Volume 1 (AMI Root) (8 GiB, EBS, General purpose SSD (gp2))' with a note: 'AMI Volumes are not included in the template unless modified'. An 'Add new volume' button is present. A warning states: 'The selected AMI contains more instance store volumes than the instance allows. Only the first 0 instance store volumes from the AMI will be accessible from the instance'.
- Instance tags**: Fields for 'Key' (set to 'Name') and 'Value' (set to 'oxclo01-') are shown. Below them is an 'Add tag' button and a note: '49 remaining (Up to 50 tags maximum)'. To the right, there are checkboxes for 'Tag instances' (checked) and 'Tag volumes' (unchecked), and a 'Remove tag' button.

44. Expand **Advanced Details** and add the same **userdata** as before.

45. Click **Create Launch Template**

46. Once you have completed the process you will be prompted to create an Auto Scaling group with this template. Click on **Create Auto Scaling Group**.

47. You will see.

1. Configure Auto Scaling group details 2. Configure scaling policies 3. Configure Notifications 4. Configure Tags 5. Review

Create Auto Scaling Group

Group name ⓘ

Launch Template ⓘ lt-0b62d8f8dfc3b7cc

Launch Template Version ⓘ Default Create new launch template

Launch Template Description ⓘ A template for a node server

Fleet Composition

☒ **Adhere to the launch template**
The launch template determines the instance type and purchase option (On-Demand or Spot).

☐ **Combine purchase options and instances**
Choose a mix of On-Demand Instances and Spot Instances and multiple instance types. Spot Instances are automatically launched at the lowest price available.

Group size ⓘ Start with instances

Network ⓘ vpc-42fb9527 (172.31.0.0/16) (default) Create new VPC

Subnet ⓘ [Create new subnet](#)

[Advanced Details](#)

48. Give the Group name as: *userid-asg* (e.g. *oxclo02-asg*).

49. Keep Group Size as 1

50. Choose **a subnet** from the options that drop down when you select that box.
Any one will do, or you can select multiple.

51. If you select the **Advanced Details** you will see that there is a “Grace Period” of 300 seconds. Read the description if you hover over the (i).

52. Click **Next: Configure Scaling Policies**

53. Leave this as **Keep this group at its original size**

Create Auto Scaling Group

You can optionally add scaling policies if you want to adjust the size (number of instances) of your group automatically. A scaling policy is a set of instructions for making such adjustments in response to an Amazon CloudWatch alarm that you assign to it. In each policy, you can choose to add or remove a specific number of instances or a percentage of the existing group size, or you can set the group to an exact size. When the alarm triggers, it will execute the policy and adjust the size of your group accordingly. [Learn more](#) about scaling policies.

- ☒ Keep this group at its initial size
- ☐ Use scaling policies to adjust the capacity of this group

54. Click **Next: Configure Notifications**

55. Click **Next: Configure Tags**

56. In the tags Key, specify **Name**, and in the Value field *userid-autoscaled* (e.g. *oxclo02-autoscaled*). Make sure that “tag new instances” is selected.

Create Auto Scaling Group

A tag consists of a case sensitive key-value pair that you can use to identify your group. For example, you could define a tag with Key = Environment and Value = Production. You can optionally choose instances in the group when they launch. [Learn more](#).

Key	Value	Tag New Instances
Name	oxclo02-autoscaled	<input checked="" type="checkbox"/>

Add tag 9 remaining

57. Click **Review**

58. Click **Create Auto Scaling Group**:

Auto Scaling group creation status



- ▼ View
- [View your Auto Scaling groups](#)
 - [View your launch configurations](#)

► Here are some helpful resources to get you started

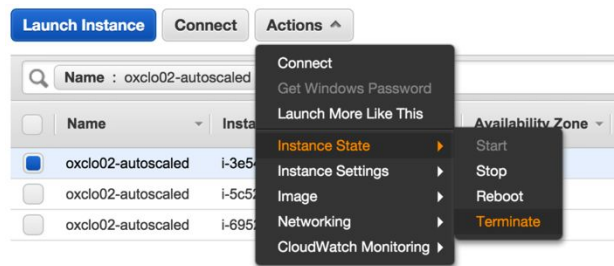
59. Now if you go to your EC2 Instances dashboard (**EC2 Dashboard -> Running Instances**), you should see a new instance starting up. Once it is started it will be tagged with your *userid-autoscaled* so you can see which ones are yours.

60. First check that your server is working properly by browsing <http://<ip-address>:8080>


You may need to be patient while the server starts up.

Don't continue to the next step until you get a proper response.

61. Now using the **EC2 Dashboard -> Running Instances** screen you can terminate this instance:



62. Click **Yes Terminate** on the next screen:



63. Now wait up (based on the previous grace period) and you should see a new instance spawned to replace the one you killed. Amazon is ensuring that you have an instance running at all times (give or take a little bit of startup time).

64. If you configured a notification, check your email.

65. Check the new server is correctly serving the data.

66. Is this newly created server much use to anyone? What would you need to do to make it more useful?
(PS I don't mean the app: let's assume that the app is in fact useful!)

67. What would you need to do to update the code on this system?

68. Go to **Auto Scaling Groups** in the left hand side of the EC2 menu. Hint it is right at the bottom!

69. You can turn off your ASG by changing the instances to 0. Click Edit on the Group Details and change the desired instances to zero (and the min/max)

Group size [X]

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity
0 [up] [down]

Minimum capacity
0 [up] [down]

Maximum capacity
0 [up] [down]

[Cancel] [Update]

70. Go and check in **Instances** that your autoscaled instance is being deleted!

71. If you just needed to stop the instance for a while, this is a simple method that doesn't involve deleting the ASG. However, we are now done, so delete the ASG!

Delete Auto Scaling groups [X]

Are you sure you want to delete this resource?

- oxclo01-asg

Deleting the Auto Scaling group will terminate all instances in the group. This action cannot be undone.

To confirm deletion, type *delete* in the field.

[Cancel] [Delete]

72. In the next lab we will reconfigure this ASG to *auto-scale* instead of having just one instance, and then test it under load.

73. **Congratulations, lab complete.**