

Exercise 8

More Apache Spark and Python, EC2

Prior Knowledge

Unix Command Line Shell

Simple Python

Learning Objectives

Using Spark on EC2

Accessing S3 files on Spark

Reading CSV files in Spark

Seeing the differences between Spark and Hadoop by performing the Word

Analysis in Spark

Spark SQL

Software Requirements

(see separate document for installation of these)

- EC2 credentials
- Flintrock

Part A. Starting Spark in EC2

1. There is a project from the creators of Spark to run it in EC2, but it is not very good! Instead we will use a tool called **flintrock**
2. Before we can use flintrock, you need to modify the config file for flintrock so that it uses your own keys. Edit the flintrock config file:

```
subl ~/.config/flintrock/config.yaml
```

It will look something like:

```
services:
  spark:
    version: 2.4.3
  hdfs:
    version: 2.7.3

provider: ec2

providers:
  ec2:
    key-name: oxclo01
    identity-file: /home/oxclo/keys/oxclo01.pem
    instance-profile-name: ec2-access-s3
    user: ec2-user
    instance-profile-name: ec2-access-s3

launch:
  num-slaves: 2
```



The source for this is here:

<https://freo.me/flintrock-conf>

This is modified in a couple of ways. Firstly, it gives the Ireland region and AMI files. Secondly, there is an “instance-profile-name”. This is a AWS feature that gives the running VM access to other APIs - in this case S3.

3. Change the key name and identity file to match your key name and identity file.
4. Make sure **install-hdfs: False**
5. Make num-slaves: **2**
6. Save the file

7. You should now be able to launch a cluster in Amazon:

```
flintrock launch oxcloXX-sc
```

(using your XX)

8. Now you should see something like:

```
Warning: Downloading Spark from an Apache mirror. Apache mirrors are often slow and unreliable, and
typically only serve the most recent releases. We strongly recommend you specify a custom download
source. For more background on this issue, please see: https://github.com/nchammas/flintrock/issues/238
Launching 3 instances...
[34.253.234.105] SSH online.
[34.253.234.105] Configuring ephemeral storage...
[52.51.185.103] SSH online.
[52.212.199.209] SSH online.
[52.51.185.103] Configuring ephemeral storage...
[52.212.199.209] Configuring ephemeral storage...
[34.253.234.105] Installing Java 1.8...
[52.51.185.103] Installing Java 1.8...
[52.212.199.209] Installing Java 1.8...
[52.51.185.103] Installing Spark...
[34.253.234.105] Installing Spark...
[52.212.199.209] Installing Spark...
[52.212.199.209] Configuring Spark master...
Spark online.
launch finished in 0:02:33.
Cluster master: ec2-52-212-199-209.eu-west-1.compute.amazonaws.com
Login with: flintrock login oxclo01-sc
```

Ignore the Apache mirror warning.

If you have issues you can try:

```
flintrock --debug launch oxcloXX-sc
```

9. Let's login to the master (all one line):

```
flintrock login oxcloXX-sc
```

You see something like:

```
Warning: Permanently added '34.253.201.139' (ECDSA) to the list of known hosts.
Last login: Mon Jul 10 18:55:35 2017 from
host109-156-251-208.range109-156.btcentralplus.com

  _| _| )
  _| ( / Amazon Linux AMI
  _|\_|_|_|

https://aws.amazon.com/amazon-linux-ami/2017.03-release-notes/
1 package(s) needed for security, out of 1 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-6-32 ~]$
```

```
pyspark --master spark://0.0.0.0:7077
--packages org.apache.hadoop:hadoop-aws:2.7.4
```

```
org.xerial.snappy@snappy-java;1.0.4.1 from central in [default]
xmlenc+xmlenc;0.52 from central in [default]
```

	modules	artifacts
conf	number	search(dwnlded evicted) number(dwnlded)
default	71 71 71 0	70 70

```
:: retrieving :: org.apache.spark#spark-submit-parent-4b5025a3-23bc-4a5f-9723-5c46c3556b4e
confs: [default]
    70 artifacts copied, 0 already retrieved (36395kB/100ms)
2018-07-19 16:16:02 WARN NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
```

Welcome to
 version 2.3.1
Using Python version 2.7.12 (default, Sep 1 2016 22:14:00)
SparkSession available as 'spark'.
>>>

14. We are going to use Spark's SQL support, which in turn uses Apache Hive.

First let's tell spark we are using SQL. In the Python command-line type:

```
from pyspark.sql import SQLContext
sqlc = SQLContext(sc)
```

```
df = sqlc.read.csv('s3a://oxclo-wind/2015/*',header='true', inferSchema='true')
```

© Paul Fremantle 2015. Licensed under the This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>



17. The df object we have is not an RDD, but instead a DataFrame. This is basically a SQL construct. (But we can easily convert it into an RDD as you will find out shortly)

18. We can print a nice table showing the first few rows with:

```
df.show(4)
```

Station_ID	Station_Name	Location_Label	Interval	MinutesInterval	End_Time	Wind_Velocity_Mtr_Sec	Wind_Direction_Variance_Deg	Wind_Direction_Deg	Ambient_Temperature_Deg	Global_Horizontal_Irradiance
1	BP (Winnerville Inlet...)	Winnerville	0.2018-01-07 00:00	1.000	0.0	100.0	0.02	0.000	0.000	0.000
1	BP (Winnerville Inlet...)	Winnerville	0.2018-01-07 00:10	1.000	0.0	101.0	0.02	0.000	0.000	0.000
1	BP (Winnerville Inlet...)	Winnerville	0.2018-01-07 00:20	1.000	0.0	102.0	0.02	0.000	0.000	0.000
1	BP (Winnerville Inlet...)	Winnerville	0.2018-01-07 00:30	1.000	0.0	103.0	0.02	0.000	0.000	0.000

(I shrunk this so you can see the table nicely!)

19. We can also convert the DataFrame into an RDD, allowing us to do functional programming on it (map/reduce/etc)

```
winds = df.rdd
```

20. Let's do the normal step of mapping the data into a simple <K,V> pair. Each column in the row can be accessed by the syntax e.g. row.Station_ID

We can therefore map our RDD with the following:

```
mapped = winds.map(lambda s: (s.Station_ID, s.Wind_Velocity_Mtr_Sec))
```

21. We can simply calculate the maximum values with this reducer:

```
maxes = mapped.reduceByKey(lambda a, b: a if (a>b) else b)
```

22. And once again collect / print:

```
for (k,v) in maxes.collect(): print k,v
```

Because python uses indentation, it can't tell if this is the end of the statement so you will see:

...

Press Enter.

23. You will see a bunch of log before the following appears:

```
SF18 10.57
SF36 11.05
SF37 7.079
SF15 7.92
SF04 34.12
SF17 5.767
```

24. You can also turn the response of a collect into a Python Map, which is handy. Try this:

```
maxes.collectAsMap()['SF04']
```

25. You can also try:
`print maxes.collectAsMap()`

PART B – Getting Jupyter running with Flintrock

26. Quit the pyspark REPL (Ctrl-D) and get back to the ec2 command line

27. Type the following commands to install and run jupyter into your master node (available here: <https://freo.me/flintrock-j>)

```
sudo yum install gcc gcc-c++ -y
sudo yum install python2.7-pip -y
sudo pip-2.7 install jupyter
export PYSARK_DRIVER_PYTHON=jupyter
export PYSARK_DRIVER_PYTHON_OPTS='notebook --no-browser'
pyspark --master spark://0.0.0.0:7077 \
--packages org.apache.hadoop:hadoop-aws:2.7.4
```

28. You will see something like:

```
[I 21:20:38.933 NotebookApp] Serving notebooks from local directory: /home/ec2-user
[I 21:20:38.934 NotebookApp] The Jupyter Notebook is running at:
[I 21:20:38.934 NotebookApp] http://localhost:8888/?token=71c8d14cb639b2c047e1e456a331b6b0e1d64f986c80370
[I 21:20:38.934 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[W 21:20:38.934 NotebookApp] No web browser found: could not locate runnable browser.
[C 21:20:38.935 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
http://localhost:8888/?token=71c8d14cb639b2c047e1e456a331b6b0e1d64f986c80370
```

29. Don't try to access that URL just yet. That is a URL that is only accessible from within the master node running on EC2 at the moment.

- 30.

To allow us to access that URL, we need to setup an SSH tunnel to the master node.

Start a new Ubuntu terminal window.

Find the name of the master node once again:

flintrock describe oxcloXX-sc

```
state: running
node-count: 3
master: ec2-34-244-248-67.eu-west-1.compute.amazonaws.com
slaves:
  - ec2-34-240-88-3.eu-west-1.compute.amazonaws.com
  - ec2-34-247-53-166.eu-west-1.compute.amazonaws.com
```

Now start ssh thus (all one line, and replace the hostname)

```
ssh -i ~/keys/oxcloXX.pem -4 -fN -L 8888:localhost:8888
ec2-user@ec2-34-244-248-67.eu-west-1.compute.amazonaws.com
```

31. Now we can open that URL in the other window. You are now accessing the Jupyter server running in EC2. Now you can use the Jupyter model as before.
32. *Note that any python code you save here will be stored on the AWS instance and deleted when you destroy the cluster!*

PART C - SQL

33. There is an easier way to do all this if you are willing to write some SQL.
34. We need to recreate the DataFrame first, so run this in a cell:

```
from pyspark.sql import SQLContext
sqlc = SQLContext(sc)
df = sqlc.read.csv('s3a://oxclo-wind/2015/*',header='true', inferSchema='true')
df.show(4)
```

35. Now we need to give our DataFrame a table name:
df.registerTempTable('wind')
36. Now we can use a simple SQL statement against our data.
ALL ON ONE Line type:

```
sqlc.sql("SELECT Station_ID, avg(Wind_Velocity_Mtr_Sec) as
avg,max(Wind_Velocity_Mtr_Sec) as max from wind group by Station_ID").show()
```

Station_ID	avg	max
SF37	2.260403505500663	7.079
SF15	1.8214145677504483	7.92
SF04	2.300981748124102	34.12
SF17	0.5183500253485376	5.767
SF18	2.2202234391695437	10.57
SF36	2.464172530911313	11.05

37. Bingo you should see a lot of log followed by:

38. Recap. So far we have:

- Started Spark in EC2
- Loaded data from S3
- Used SQL to read in CSV files
- Explored Map/Reduce on those CSV files
- Used SQL to query the data.

39.

Find the IP address of the Spark Master: in your Ubuntu start a new terminal and type:
flintrock describe oxclo01-sc


You should see something like:

```
oxclo01-sc:
state: running
node-count: 3
master: ec2-52-214-61-215.eu-west-1.compute.amazonaws.com
slaves:
- ec2-34-240-42-233.eu-west-1.compute.amazonaws.com
- ec2-34-245-14-42.eu-west-1.compute.amazonaws.com
```

40. Go to e.g.

<http://ec2-52-214-61-215.eu-west-1.compute.amazonaws.com> :8080 using the master's DNS address (not the one in this text)

You should see something like:

 **Spark Master at [spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077](http://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077)**

URL: [spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077](http://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077)
REST URL: [spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:6066](http://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:6066) (cluster mode)

Alive Workers: 1
Cores in use: 1 Total, 1 Used
Memory in use: 2.7 GB Total, 1024.0 MB Used
Applications: 1 [Running](#), 0 [Completed](#)
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers

Worker Id	Address	State	Cores	Memory
worker-20170710185543-172.31.1.109-39733	172.31.1.109:39733	ALIVE	1 (1 Used)	2.7 GB (1024.0 MB Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20170710192533-0000	(kill) PySparkShell	1	1024.0 MB	2017/07/10 19:25:33	ec2-user	RUNNING	10 s

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

41. If you want you can try adding another slave and then rerun the analysis.

You can see the extra core working in the Web UI

flintrock add-slaves --num-slaves 1 oxcloXX-sc

If you need it the code is here:

<https://freo.me/wind-sql>

42.

We must remember to stop our cluster as well (its costing money...)
From Ubuntu terminal

```
flintrock destroy oxcloXX-sc
```

Type y when prompted.

43. Congratulations, this lab is complete.